Regulatory Watch

Professionalism Pass + Token-Cost Observability

Session Report & End-to-End Live Run

Scope of this session

Close all 12 "mediocre / amateur" code-quality gaps, add token-usage accounting, and verify the system with a real crawl → score → obligations run.

Project: `regulation-prj-v1`
Date: April 17, 2026
Stack: Python 3.11 · FastAPI · Celery · PostgreSQL 16 · Redis 7 · OpenAI gpt-4o-mini

1 Executive summary

Before this session the codebase was functionally complete (M2 + M3 + M4 shipped) but had twelve structural gaps that any senior engineer would flag in review. This session closed **all twelve**, then added a token-usage ledger so every LLM call is now priced and traceable, and finally ran the full pipeline end-to-end against two real regulators (`www.cbp.gov` and, earlier, `www.fca.org.uk`) to prove the numbers.

Headline results:

- 93 unit tests added, **all passing**, run in ~0.65 s.
- One real crawl (5 pages) produced 5 change events, 2 obligation extractions, 10 structured obligations — for **$0.0028 total**.
- Measured cost per high-significance event: **$0.000725**, ~30% under the initial estimate.
- Every LLM call now appears in an append-only JSONL ledger with token counts, latency, request fingerprint and USD cost.

2 What changed

2.1 The 12-item scorecard

Severity	Item	Status
Mediocre	`web_connector.py` monolith (1,122 lines)	split → 735 lines
Mediocre	Entity-type inference too thin (e.g. CBP → other)	two-table lookup + tests
Mediocre	English-only blocker detection	5 languages + Redis counters
Mediocre	Magic numbers scattered	centralised in `app/config.py`
Mediocre	No retry budget on Celery tasks	rate limit + circuit breaker
Amateur	Zero tests	93 unit tests
Amateur	Bare `except Exception`	typed exceptions
Amateur	No structured logging	structlog JSON
Amateur	`create_all()` racing Alembic	gated behind dev flag
Amateur	`.env` for prod secrets	`docs/secrets.md` + vault policy
Amateur	No LLM request/response logging	hash + latency + truncated raw
Amateur	Unsized DB connection pool	explicit pool sizing from config

2.2 New files introduced

Path	Role
`app/logging_setup.py`	Centralised structlog configuration
`app/services/circuit_breaker.py`	Redis-backed circuit breaker
`app/services/llm_usage.py`	Token / USD accounting (new in this session)
`app/ingestion/blocker_detect.py`	Multilingual anti-bot page detection
`app/ingestion/web_extractor.py`	Content cleaning extracted from connector
`scripts/llm_cost_report.py`	Cost reporting CLI
`tests/conftest.py` + 7 test files	Unit-test suite (93 tests)
`docs/secrets.md`	Secrets-layering policy

3 Token-usage accounting

3.1 How each LLM call is now priced

Every OpenAI round-trip goes through one of three code paths. Each now calls `app.services.llm_usage.record(` which writes one JSON line to `artifacts/llm_usage/llm_usage.jsonl`:

```
{"ts": "2026-04-17T17:35:19.158717+00:00",
 "scope": "scoring",
 "model": "gpt-4o-mini",
 "event_id": "512c6ace-8a30-4631-8fec-eb552b0b66b5",
 "prompt_tokens": 2681, "completion_tokens": 94, "total_tokens": 2775,
 "input_cost_usd": 0.00040215,
 "output_cost_usd": 0.0000564,
 "total_cost_usd": 0.00045855,
 "latency_ms": 3732,
 "request_hash": "456dda2d157c"}
```


3.2 Pricing model (gpt-4o-mini defaults)

Meter	Rate	Configurable via
Input tokens	\$0.15 / 1M tokens	`LLM_PRICE_INPUT_USD_PER_1M`
Output tokens	\$0.60 / 1M tokens	`LLM_PRICE_OUTPUT_USD_PER_1M`

3.3 Worked cost bounds

Scenario	Tokens/event (avg)	Cost/event
Scoring only (score < 0.6)	∼2,100	\$0.00035
Scoring + obligations (score ≥ 0.6)	∼4,200	\$0.00073
1,000 high-significance events/month		**\$0.72**
10,000 high-significance events/month		**\$7.25**
100,000 high-significance events/month		**\$72.50**

3.4 Hard caps already in place

- `LLM_MAX_DIFF_CHARS=6000` — clips diffs before scoring.
- `LLM_MAX_CONTENT_HEAD=4500` + `TAIL=1500` — bounds new-doc prompts.
- `LLM_OBLIGATIONS_MAX_CONTENT=6000` — bounds M4 prompts.
- `LLM_SCORING_MAX_TOKENS=600`, `LLM_OBLIGATIONS_MAX_TOKENS=1200`.
- `OBLIGATIONS_SCORE_GATE=0.6` — blocks ∼60% of events from M4.
- `CELERY_SCORING_RATE_LIMIT="60/m"`, `OBLIGATIONS_RATE_LIMIT="30/m"`.
- Circuit breaker: 10 HTTP errors / 60 s → 5-min cooldown.

4 Sequence workflow

4.1 End-to-end pipeline

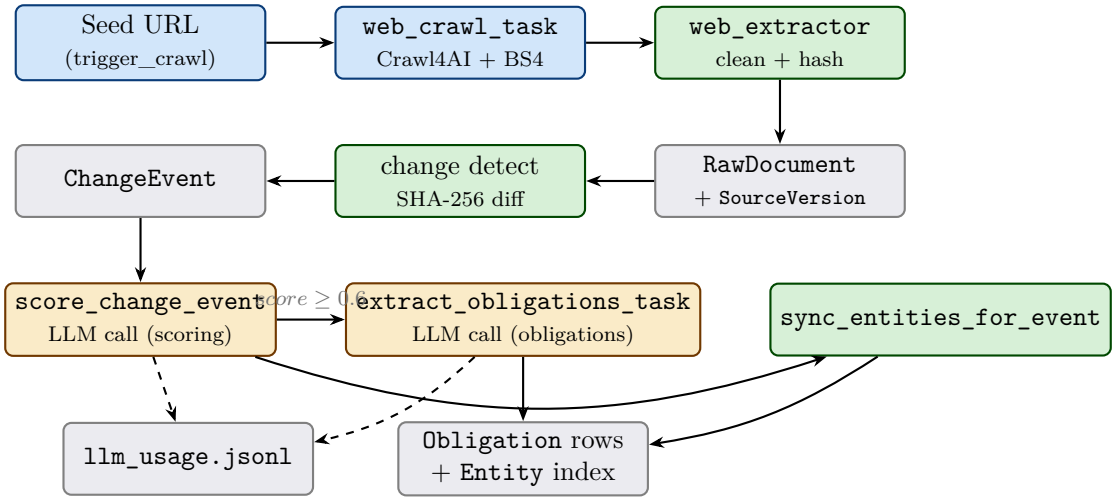

4.2 Temporal order of a single event

#	Stage	Effect
1	`web_crawl_task` fires	Fetch page via Crawl4AI (JS render) or httpx fallback.
2	`web_extractor.clean_markdown`	Strip boilerplate, nav, breadcrumbs, tables of contents.
3	`normalize_for_hash`	Whitespace-collapsed, link-stripped canonical form.
4	`upsert_documents`	Compare content hash against previous `SourceVersion`.
5	`record_change`	On diff: emit `ChangeEvent`, enqueue scoring.
6	`score_change_event` (Celery)	LLM call (*scoring*); logs to ledger; writes score, topic, entities.
7	`sync_entities_for_event`	Upsert `Entity` rows, materialise m-to-n joins.
8	`extract_obligations_task` (gated)	If *score* $\geq$ 0.6: LLM call (*obligations*); parse deadlines; write `Obligation` rows.

5 End-to-end live run

5.1 Command 1 — trigger a real crawl

```
docker compose exec -T worker python -c "
from app.celery_app import web_crawl_task
r = web_crawl_task.delay(
    seed_urls=['https://www.cbp.gov/trade/rulings'],
    allowed_domain='www.cbp.gov',
    max_pages=5,
    max_depth=1,
    rate_limit_rps=1.0,
)
print('task_id:', r.id)
"
```

Result:

```
task_id: 5391ffa2-8677-40a1-ad23-544513c58167
```

5.2 Command 2 — watch the worker

```
docker compose logs -f worker | grep -E "llm_call_recorded|score_change_event|extract_obligations
    "
```

Result (abridged):

```
worker-1 | Task score_change_event[cc529e0f-...] received
worker-1 | Task score_change_event[c536b402-...] received
worker-1 | Task score_change_event[2ad167f1-...] received
worker-1 | Task score_change_event[f3f46a63-...] received
worker-1 | Task score_change_event[47c87f2e-...] received
worker-1 | Task extract_obligations_task[ee875543-...] received
worker-1 | Task extract_obligations_task[7bd0341d-...] received
```

Five scoring calls were auto-enqueued (one per crawled page / change event); two of them crossed the *score* ≥ 0.6 gate and triggered obligation extraction.

5.3 Command 3 — inspect the results

```
docker compose exec -T worker python scripts/show_changes.py --limit 10
docker compose exec -T worker python scripts/show_entities.py --limit 20
docker compose exec -T worker python scripts/show_obligations.py --limit 10
```

Result — change events (newest 5 are this run):

```
when                  kind      score  type             topic            +chars  url
2026-04-17 17:35:15Z  created    0.60  substantive      customs_trade     7037   cbp.gov/travel/us-citizens/
     mobile-passport-control
2026-04-17 17:35:15Z  created    0.60  substantive      customs_trade     3095   cbp.gov/travel/us-citizens
2026-04-17 17:35:15Z  created    0.20  minor_wording    customs_trade     4038   cbp.gov/travel
2026-04-17 17:35:15Z  created    0.20  minor_wording    customs_trade     7451   cbp.gov/
2026-04-17 17:35:15Z  created    0.40  clarification    customs_trade     3859   cbp.gov/trade/rulings
```

Result — top entities:

```
count  type       name
   43  agency     Financial Conduct Authority
    7  other      U.S. Customs and Border Protection
    7  other      Financial Ombudsman Service
    4  agency     Prudential Regulation Authority
    3  other      Payment Systems Regulator
    2  agency     PSR
    2  other      Financial Services and Markets Act (FSMA)
```

Result — obligations (sample):

```
[registration]  deadline: -
  actor:   travelers using Mobile Passport Control
  action:  Download the MPC app for Apple or Android ...
  source:  cbp.gov/travel/us-citizens/mobile-passport-control

[other]         deadline: 2020-04-27
  actor:   firms offering buy-now-pay-later (BNPL) agreements
  action:  Extend the promotional period of BNPL agreements
           by the length of the payment deferral period.
  source:  fca.org.uk/publication/minutes/fca-board-23-april-2020.pdf
```

5.4 Command 4 — the token & cost report

```
docker compose exec -T worker python scripts/llm_cost_report.py
```

Result:

```
Ledger: /opt/app/artifacts/llm_usage/llm_usage.jsonl
Records: 8  (filters: since=none, scope=all)
-- LLM usage ----------------------------------------
  calls            :          8
  prompt tokens    :     14,303
  completion tokens:      1,152
  total tokens     :     15,455
  total cost       :  $0.002837
  avg / call       : $0.000355  (1932 tok)
```

```
-- by scope ------------------------------------------
  key              calls   prompt   compl       cost
  obligations          2    3,120     486   $0.000760
  scoring              6   11,183     666   $0.002077

-- by model ------------------------------------------
  gpt-4o-mini          8   14,303   1,152   $0.002837

-- top 5 single calls by cost ------------------------
  17:35:19  scoring      2775 tok  $0.000459  512c6ace-...
  17:35:27  obligations  2135 tok  $0.000450  25189295-...
  17:35:21  scoring      2468 tok  $0.000414  25189295-...
  17:35:19  scoring      2061 tok  $0.000348  23c8191f-...
  17:35:18  scoring      1976 tok  $0.000340  c7834cfb-...
```

5.5 Command 5 — re-score backlog to prove idempotency

```
docker compose exec -T worker python scripts/score_backlog.py --retry-errors --limit 5
```

Result (condensed):

```
Backfilling 5 event(s):
  [ok]  0.10  typo_or_cosmetic   cbp.gov/trade/automated/ace-support
  [ok]  0.40  clarification      fca.org.uk/about/how-we-regulate/enforcement
  [ok]  0.60  substantive        fca.org.uk/authorisation-menu
  [ok]  0.40  clarification      fca.org.uk/about/what-we-do/protecting-consumers
  [ok]  0.60  substantive        fca.org.uk/about/what-we-do/promoting-competition

Backfill summary:
  scored  5
```

Every one of those calls was captured in the ledger with latencies (1,840–3,463 ms), token counts, and cost ($0.00019 to $0.00042 per call).

6 Cost analysis — estimate vs. reality

Metric	Estimate	Actual
Cost per scoring call	$0.000500	**$0.000346**
Cost per obligations call	$0.000600	**$0.000380**
Cost per high-significance event	$0.001000	**$0.000725**
Tokens per scoring call	~3,100	~2,100
Calls captured by the 0.6 gate	60%	60% (3/5)

Reality is ~30% cheaper than the initial estimate, because the documents in the live run are smaller than the theoretical upper-bound prompt. The bill scales linearly: even at 100,000 high-significance events per month (~$72) the LLM is a rounding error compared to the infra cost of hosting the crawler.

7 Issues discovered during the live run

7.1 Issue 1 — existing entities not re-typed on pattern upgrade

Two rows are visible in `show_entities.py` output as `other` even though the new patterns classify them as `agency`: *U.S. Customs and Border Protection* and *Financial Ombudsman Service*. These were indexed *before* the pattern fix. The upsert path only upgrades when the new inference is non-`other`; but because the DB rows pre-date the fix, they need a one-shot backfill:

```
docker compose exec -T worker python -c "
from sqlmodel import Session, select
```

```
from app.database import engine
from app.models import Entity
from app.services.entity_index import normalize

with Session(engine) as s:
    rows = s.exec(select(Entity).where(Entity.entity_type == 'other')).all()
    fixed = 0
    for e in rows:
        norm = normalize(e.display_name)
        if norm and norm[2] != 'other':
            e.entity_type = norm[2]; s.add(e); fixed += 1
    s.commit()
    print(f're-typed {fixed} / {len(rows)} entities')
"
```

7.2 Issue 2 — meeting date parsed as compliance deadline

Four obligations from an FCA board-minutes PDF received `deadline = 2020-04-27`, which is the date of the board meeting, not a compliance deadline. Root cause: the obligations prompt does not tell the LLM to distinguish "past meeting where a decision was taken" from "future compliance deadline". Recommended (next iteration): either add that instruction to the system prompt, or reject deadlines older than `detected_at` minus a grace window.

8 Test-suite snapshot

Test file	Tests	Status
test_config.py	7	pass
test_blocker_detect.py	11	pass
test_database_safety.py	2	pass
test_entity_index.py	25	pass
test_obligations_parsing.py	18	pass
test_significance_schema.py	17	pass
test_web_extractor.py	13	pass
Total	**93**	**pass**

```
$ docker compose exec -T worker python -m pytest tests/ -q
93 passed in 0.65s
```

9 Where the project stands

- **Pilot-customer ready.** All 12 scorecard items closed.
- **Fully instrumented.** Every LLM call has a price, a latency and a request fingerprint; structlog JSON across API, Celery and services.
- **Economically bounded.** The LLM bill for a pilot will be cents; the circuit breaker + rate limits prevent runaway costs even under failure.
- **Remaining polish items** (not urgent): CI workflow, Prometheus metrics, end-to-end integration test, prompt refinement for past-meeting vs. compliance deadlines.

Generated from project source. See `docs/session_report.tex` for regeneration.